



Searching with ElasticSearch

Richard Miller
Jeremy Mikola

Symfony Live London
September, 14, 2012

Who are we?



Richard Miller, [@mr_r_miller](#)
Software Engineer [@SensioLabsUK](#)



Jeremy Mikola, [@jmikola](#)
Software Engineer [@10gen](#)

Search is important.

**Database searching is good enough,
though, right?**

But search engines are difficult to install, configure, and use.

ElasticSearch

<http://www.elasticsearch.org/>

ElasticSearch in a nutshell

- Based on Lucene
- Schema-less
- RESTful
- Document-oriented (JSON)
- Fast and scalable

Who is using ElasticSearch?

- Mozilla
- StumbleUpon
- Klout
- GOV.UK
- FourSquare

Getting Started

- Download
 - <http://www.elasticsearch.org/download/>
- Launch
 - Shell script (background/foreground)
 - Service
- Configuration (optional)
 - Runtime parameters
 - File-based
 - REST API

Indexes, Types

URL Structure

<http://localhost:9200/cookbook/recipes/>

Inserting Documents

```
$ curl -XPOST http://localhost:9200/cookbook/recipes -d '{  
  "name": "Welsh Rarebit",  
  "tags": ["cheese", "bread"]  
'  
  
{  
  "ok": true,  
  "_index": "cookbook",  
  "_type": "recipes",  
  "_id": "IcYOL_NuT-ymRwl4Iz2NyA",  
  "_version": 1  
}
```

Inserting Documents

```
$ curl -XPOST http://localhost:9200/cookbook/recipes/2 -d '{  
  "name": "Beef Wellington",  
  "tags": ["beef"]  
'  
  
{  
  "ok": true,  
  "_index": "cookbook",  
  "_type": "recipes",  
  "_id": "2",  
  "_version": 1  
}
```

Inserting Documents

```
$ curl -XPUT http://localhost:9200/cookbook/recipes/3 -d '{  
  "name": "Yorkshire Pudding",  
  "tags": ["pastry"]  
'  
  
{  
  "ok": true,  
  "_index": "cookbook",  
  "_type": "recipes",  
  "_id": "3",  
  "_version": 1  
}
```

Updating Documents

```
$ curl -XPOST http://localhost:9200/cookbook/recipes/2 -d '{  
  "name": "Beef Wellington",  
  "tags": ["beef", "steak", "pastry"]  
}'  
  
{  
  "ok": true,  
  "_index": "cookbook",  
  "_type": "recipes",  
  "_id": "2",  
  "_version": 2  
}
```

Basic Searching with URI Requests

```
$ curl -XGET http://localhost:9200/cookbook/recipes/_search?q=tags:pastry
```

```
{
  "took": 31, "timed_out": false, "_shards": { "total": 5, "successful": 5, "failed": 0
},
  "hits": {
    "total": 2, "max_score": 0.5, "hits": [
      { "_index": "cookbook", "_type": "recipes", "_id": "2", "_score": 0.5,
        "_source" : { "name": "Beef Wellington", "tags": ["beef", "steak", "pastry"]} }
    },
      { "_index": "cookbook", "_type": "recipes", "_id": "3", "_score": 0.30685282,
        "_source" : { "name": "Yorkshire Pudding", "tags": ["pastry"]} }
    ]
  }
}
```


Querying Across Indexes and Types

```
$ curl -XGET http://localhost:9200/cookbook/recipes,foods/_search?q=tags:pastry
```

```
$ curl -XGET http://localhost:9200/cookbook/_search?q=tags:pastry
```

```
$ curl -XGET http://localhost:9200/cookbook,guide/_search?q=tags:pastry
```

```
$ curl -XGET http://localhost:9200/_all/recipes/_search?q=tags:pastry
```

```
$ curl -XGET http://localhost:9200/_search?q=tags:pastry
```

Advanced Searching with Query DSL

- Basic queries
 - Term(s)
 - Prefix
 - Fuzzy
 - Range
- Compound queries
 - Bool
 - Disjunction max
 - Constant score
- Filtered
- Faceted
- "More like this"

Filters

- Not scored
- Cacheable
- Familiar operators
- Boolean logic
- Geospatial

Query DSL in JSON Request Body

```
$ curl -XGET http://localhost:9200/cookbook/recipes/_search -d '{
  "query": { "fuzzy": { "name": "Wellington" } },
  "filter": { "term": { "tags": "beef" } }
}'


{
  "took": 5, "timed_out": false, "_shards": { "total": 5, "successful": 5, "failed": 0 },
  "hits": {
    "total": 1, "max_score": 0.625, "hits": [
      { "_index": "cookbook", "_type": "recipes", "_id": "2", "_score": 0.625,
        "_source": { "name": "Beef Wellington", "tags": ["beef", "steak", "pastry"] } }
    ]
  }
}
```

Mappings

- Associated with types
- Dynamic (schema-less) by default
- Field types
 - Core
 - Objects
 - Arrays
 - Special (IP, geo, files)
- Analyzers
 - Defined at index level, assigned to types and fields
 - Stopwords, n-grams, stemming

Distributed Architecture

- Sharding
- Replication
- Node discovery
- Scatter/gather search
- Request redirection
- Automatic balancing, failover
- Multi-tenant indexes



Additional Features

- Update API
- Routing
- Parents and children
- Timestamps
- TTL

Elastica

<https://github.com/rufin/Elastica/>

Elastica is a PHP client for Elasticsearch.

```
$elasticaClient = new Elastica_Client();
```

```
$elasticaClient = new Elastica_Client([  
    'host' => 'mydomain.org',  
    'port' => 12345,  
]);
```

```
$elasticaClient = new Elastica_Client([  
    'servers' => [  
        [ 'host' => 'localhost', 'port' => 9200 ],  
        [ 'host' => 'localhost', 'port' => 9201 ],  
    ]  
]);
```

```
$elasticaIndex = $elasticaClient->getIndex('cookbook');  
$elasticaType = $elasticaIndex->getType('recipes');
```

Searching is straightforward.

```
$elasticaClient = new Elastica_Client();  
$elasticaIndex = $elasticaClient->getIndex('cookbook');  
$elasticaType = $elasticaIndex->getType('recipes');  
  
$resultSet = $elasticaType->search('pastry');
```

Search results have an object-oriented representation, too.

```
/** @var Elastica_ResultSet */  
$resultSet = $elasticaType->search('pastry');  
  
$totalHits = $resultSet->getTotalHits();  
  
foreach ($resultSet->getResults() as $result) {  
    /** @var Elastica_Result */  
    $result->getScore();  
    $result->getExplanation();  
    $data = $result->getData();  
}
```

Complex searches may be built up from objects.

```
$mltQuery = new Elastica_Query_MoreLikeThis();  
$mltQuery->setLikeText('a sample recipe');  
$mltQuery->setFields(['name', 'description']);
```

```
$existsFilter = new Elastica_Filter_Exists('ingredients');  
$notTagFilter = new Elastica_Filter_Not(  
    new Elastica_Filter_Term([ 'tags', 'stodgy' ])  
);
```

```
$andFilter = new Elastica_Filter_And();  
$andFilter->addFilter($notTagFilter);  
$andFilter->addFilter($existsFilter);
```

```
$mltQuery->setFilter($andFilter);
```

Mappings may be set for types.

```
$elasticaType = $elasticaIndex->getType('recipes');
```

```
$mapping = new Elastica_Type_Mapping($elasticaType);  
$mapping->setProperties([  
    'name' => [ 'type' => 'string', 'boost' => 5 ],  
    'tags' => [ 'type' => 'string', 'index_name' => 'tag', 'boost' => 3 ],  
]);  
$mapping->send();
```

// Alternatively...

```
$elasticaType->setMapping([  
    'name' => [ 'type' => 'string', 'boost' => 5 ],  
    'tags' => [ 'type' => 'string', 'index_name' => 'tag', 'boost' => 3 ],  
]);
```

It'd be helpful if searches matched more than just complete words.

```
$elasticsearch->create([  
  'number_of_shards' => 4,  
  'number_of_replicas' => 2,  
  'analysis' => [  
    'analyzer' => [  
      'indexAnalyzer' => [ 'type' => 'snowball' ],  
      'searchAnalyzer' => [ 'type' => 'snowball' ],  
    ],  
  ],  
], true);
```

// If we search for "baking", we can get results for "baked", "bakes", etc.

Analyzers may be applied to types.

```
$mapping = new Elastica_Type_Mapping($elasticaType);  
$mapping->setParam('index_analyzer', 'indexAnalyzer');  
$mapping->setParam('search_analyzer', 'searchAnalyzer');  
$mapping->setProperties(  
    // ...  
);  
$mapping->send();
```

We also want to avoid hits for uninteresting, common words.

```
$elasticaIndex->create([
  'analysis' => [
    'analyzer' => [
      'url_analyzer' => [
        'type' => 'custom',
        'tokenizer' => 'lowercase',
        'filter' => [ 'stop', 'url_stop' ],
      ],
    ],
    'filter' => [
      'url_stop' => [ 'type' => 'stop', 'stopwords' => [ 'http', 'https' ] ],
    ],
  ], true);
```


Analyzers may also apply to specific fields.

```
$mapping->setProperties([  
    'url' => [ 'type' => 'string', 'analyzer' => 'url_analyzer' ],  
    // ...  
]);  
  
$urlQuery = new \Elasticsearch\Query_Text();  
$urlQuery->setFieldQuery('url', 'pastry');  
$urlQuery->setFieldParam('url', 'analyzer', 'url_analyzer');
```

We may want to present faceted navigation for a search.

```
$elasticaFacet = new Elastica_Facet_Terms('myFacetName');  
$elasticaFacet->setField('tags');  
$elasticaFacet->setSize(10);  
  
// Add that facet to the search query object.  
$elasticaQuery->addFacet($elasticaFacet);
```

Facet data will be included with query results.

```
// Get facets from the result of the search query
```

```
$elasticaFacets = $elasticaResultSet->getFacets();
```

```
// Note: "myFacetName" is the name of the facet we defined
```

```
foreach ($elasticaFacets['myFacetName']['terms'] as $elasticaFacet) {  
    printf("%s: %s\n", $elasticaFacet['term'], $elasticaFacet['count']);  
}
```

```
beef: 3
```

```
pastry: 2
```

```
roast: 1
```

```
pie: 1
```

FOQElasticaBundle

<https://github.com/Exercise/FOQElasticaBundle>

Indexes and types are defined in the bundle's configuration.

foq_elastica:

clients: { default: { host: localhost, port: 9200 } }

indexes:

cookbook:

client: default

types:

recipes: ~

Mapping Types to DB Entities

```
# foq_elastica / indexes / types
```

```
  recipes:
```

```
    mappings:
```

```
      name: { type: string, boost: 5 }
```

```
      tags: { type: string, boost: 3 }
```

```
    persistence:
```

```
      driver: orm
```

```
      model: CookbookBundle\Entity\Recipe
```

```
      provider: { query_builder_method: createIsPublishedQueryBuilder }
```

```
      listener: { is_indexable_callback: isPublished }
```

Index and Type Services

```
/** @var Elastica_Index */
```

```
$cookbookIndex = $this->container->get('foq_elastica.index.cookbook');
```

```
/** @var Elastica_ResultSet */
```

```
$resultSet = $cookbookIndex->search('pastry');
```

```
/** @var Elastica_Type */
```

```
$recipesType = $this->container->get('foq_elastica.index.cookbook.recipes');
```

```
/** @var Elastica_ResultSet */
```

```
$resultSet = $recipesType->search('pastry');
```

Transforming Search Results

```
# foq_elastica / indexes / types
```

```
  recipes:
```

```
    persistence:
```

```
      driver: orm
```

```
      model: CookbookBundle\Entity\Recipe
```

```
      finder: ~
```

```
/** @var FOQ\ElasticaBundle\Finder\TransformedFinder */
```

```
$finder = $container->get('foq_elastica.finder.cookbook.recipes');
```

```
/** @var array of CookbookBundle\Entity\Recipe objects */
```

```
$recipes = $finder->find('pastry');
```


Results and Entities Together

```
/** @var array of FOQ\ElasticaBundle\HybridResult */
$hybridResults = $finder->findHybrid('pastry');

foreach ($hybridResults as $hybridResult) {
    /** @var CookbookBundle\Entity\Recipe */
    $recipe = $hybridResult->getTransformed();

    /** @var Elastica_Result */
    $elasticaResult = $hybridResult->getResult();
}
```

Console Commands

```
$ php app/console foq:elastica:populate --index cookbook --no-debug
```

```
$ php app/console foq:elastica:search --index cookbook --type recipes \  
  --query pastry --show-field name
```

Repository Classes

foq_elastica / indexes / types

recipes:

 persistence:

 driver: orm

 model: CookbookBundle\Entity\Recipe

 repository: CookbookBundle\Search\RecipeRepository

Complex queries can be encapsulated in repositories.

```
<?php
```

```
use FOQ\ElasticaBundle\Repository;
```

```
namespace CookbookBundle\Search;
```

```
class RecipeRepository extends Repository
```

```
{
```

```
    public function findWithCustomQuery($searchText)
```

```
    {
```

```
        // Build complex $query with Elastica objects
```

```
        return $this->find($query);
```

```
    }
```

```
}
```

Querying with Repository Services

```
/** @var FOQ\ElasticaBundle\Manager\RepositoryManager */  
$repositoryManager = $container->get('foq_elastica.manager');
```

```
/** @var FOQ\ElasticaBundle\Repository */  
$repository = $repositoryManager->getRepository(CookbookBundle:Recipe');
```

```
/** @var array of CookbookBundle\Entity\Recipe */  
$recipes = $repository->findWithCustomQuery('pastry');
```

Indexing Files

```
# foq_elastica / indexes / types
```

```
  recipes:
```

```
    mappings:
```

```
      attachedFile: { type: attachment }
```

WDT and Profiler Integration



Today

Journal

Progress

Calorie Log

Leaderboard

Photos

FAQs

I will Lose 25 lbs -
fixture goal update
a few minutes ago

Friends



Add / Invite Friends

Start a Program Today!

View our programs and choose the program that's right for you.

[View All Programs](#)

Quick Journal

How are you feeling?

Your journal may be visible to the public.

[Review your privacy settings](#)

[Publish Entry](#)

Feed: **Everyone**



You set a new goal

I will Lose 25 lbs - fixture goal update

a few minutes ago • [Add a Comment](#)



Ted left you a note

This is a note left in the UserBundle fixtures file.

a few minutes ago • [Add a Comment](#) • [View Ted's Profile](#)

Calories Allowed

[View All](#)

[Calculate](#)

+ Activity

+ Food

Weight

[+ Update](#)

Weight in Lbs

Date

Today



[Track](#)

9/9 9/10 9/11 9/12 9/13 9/14

[Add/Remove Tracking](#)

Complete Orientation

Do each item below to complete your profile

[Set a goal weight](#)

[Post a journal entry](#)

[Log a food item](#)



2.1.0-DEV

PHP 5.4.6-2~precise+1 | [xdebug](#) | [accele](#)



exercise|dev|debug|5052e5f990d13



UserController::homeAction|fos_user_user_home|200



0 ms



46.5 MB



fred









77



20



 CONFIG REQUEST EXCEPTION EVENTS LOGS TIMELINE ROUTING SECURITY DOCTRINE MONGODB 77 ELASTICA 20 SEARCHProfile for: GET <http://exercise.local/home> by 127.0.0.1 at Fri, 14 Sep 2012 04:08:25 -0400

Queries

Path: exercise/user/5052e5d5e84df1bb5c000028
Method: DELETE
{ }
Time: 3.90 ms

Path: exercise/user/5052e5d5e84df1bb5c000028
Method: PUT
{ username: fred, firstName: null, lastName: null, aboutMe: null, userStatus: null, displayLocation: null, occupation: null }
Time: 1.22 ms

Path: exercise/all_users/5052e5d5e84df1bb5c000028
Method: DELETE
{ }
Time: 1.11 ms

Path: exercise/all_users/5052e5d5e84df1bb5c000028
Method: PUT
{ username: fred, email: fred@site.org, firstName: null, lastName: null, aboutMe: null, userStatus: null, displayLocation: null, occupation: null }
Time: 1.22 ms

Path: exercise/user/5052e5d5e84df1bb5c000028
Method: DELETE
{ }
Time: 1.16 ms

Path: exercise/user/5052e5d5e84df1bb5c000028
Method: PUT
{ username: fred, firstName: null, lastName: null, aboutMe: null, userStatus: null, displayLocation: null, occupation: null }
Time: 1.49 ms

Path: exercise/all_users/5052e5d5e84df1bb5c000028
Method: DELETE

Future Plans

- Annotation-based mapping
 - Methods and/or properties
- Improve ORM/ODM agnosticity
 - Propel support is incomplete
- Elasticsearch ODM?
 - Elasticsearch is a document store
 - Is transforming to DB entities always necessary?

Final Takeaways

- Applications benefit from *real* search
- Elasticsearch is one answer
 - Simple to get up and running
 - Depth of functionality
- FOQElasticaBundle can help
 - Elastica is well-designed
 - Integrates with services and ORM/ODM
- You can improve searching in your app today

Thanks!

Questions?

<https://joind.in/7058>